



Automatic Snowfall Detection

Machine Vision for Identification of Snowfall in
Images using a Convolutional Neural Network



The Problem

Can a neural network be trained to automatically recognize falling snow via machine vision?

**U.S. Department of Transportation:
~ 4% of vehicle crashes are related to snow**

**Localized snow detection could be useful for
transportation and emergency personnel.**

Navigation apps could inform users of falling snow.

**On a separate note, ski resorts could easily monitor
new snowfall on each run.**

Presentation Overview



01.

The Data

- Data source
- Overview

02.

Preprocessing

- Preparing the data for modelling

03.

The Models

- Structure
- Performance
- Streamlit app

04.

Conclusions

- Takeaways
- Future Directions

The Data

DesnowNet: Context-Aware Deep Network for Snow Removal

Liu, Yun-Fu and Jaw, Da-Wei and Huang, Shih-Chia and Hwang, Jenq-Neng

- 50,000 images without snow
- A copy of each image with falling snow artificially added
- All RGB-encoded .jpg
- Various resolutions
- Largest dimension for any image = 640 pixels

Kindly provided for download on their website:

<https://sites.google.com/view/yunfuliu/desnownet>





Challenges of this Dataset

1. 100,000 images is far too large for working memory
 - a. How can we shuffle and split?
 - b. How can we feed so many images to the model?
2. Images must be paired with a numerical label
 - a. 1 for snow
 - b. 0 for clear
3. Long time to train
4. Images are different sizes

Preparing the Images for Modelling

1. 100,000 images is far too large for working memory
 - a. Start with just the file paths
 - b. Shuffle and split before reading in image data
 - c. Read in the images in batches

2. Images must be paired with a numerical label

3. Long time to train

Use a GPU →



4. Images are different sizes

Get the labels from the folder name



→ 0



→ 1

Buffer →



Ready for the Neural Net



1 for snow
0 for clear

Buffers make every
image 640x640

The Models



Prediction



Convolutional Layer: Groups pixels near each other, helping identify shapes.



Max Pooling Layer: Pixels near each other all given the value of the largest.



Dense Layer: Look for patterns. Mostly a black box.


Model 1

Convolution Kernel: 5x5


Max Pooling Kernel: 3x3

Training Passes: 5


Training and Validation Accuracy at the End of Each Training Epoch



Model 1



Convolution Kernel: 5x5



Max Pooling Kernel: 3x3

Training Passes: 5

Metrics

Precision: If I say a picture has snow, how often am I right?

- 99.26%

Recall: Of the pictures with snow, how many did I identify?

- 97.80%

Accuracy: Overall, how many of my predictions were correct?

- 98.53%

I don't want to miss snowy pictures.

Can I improve the Recall?


And keep high accuracy?

Model 1


Convolution Kernel: 5x5

Max Pooling Kernel: 3x3


Training Passes: 5



Model 1



Convolution Kernel: 5x5



Max Pooling Kernel: 3x3

Training Passes: 5

Updated Metrics

Precision: If I say a picture has snow, how often am I right?

- 99.26% → 98.73%

Recall: Of the pictures with snow, how many did I identify?

- 97.80% → 98.84%

Accuracy: Overall, how many of my predictions were correct?

- 98.53% → 98.78%

Cutoff can be adjusted
significantly without much
accuracy loss.

Model is highly tunable.


Model 2

Convolution Kernel: 4x4


Max Pooling Kernel: 4x4

Training Passes: 10


Training and Validation Accuracy at the End of Each Training Epoch



Model 2



Convolution Kernel: 4x4



Max Pooling Kernel: 4x4

Training Passes: 10

Metrics

Precision: If I say a picture has snow, how often am I right?

- 99.37%

Recall: Of the pictures with snow, how many did I identify?

- 99.57%

Accuracy: Overall, how many of my predictions were correct?

- 99.47%

This model is already very
balanced.

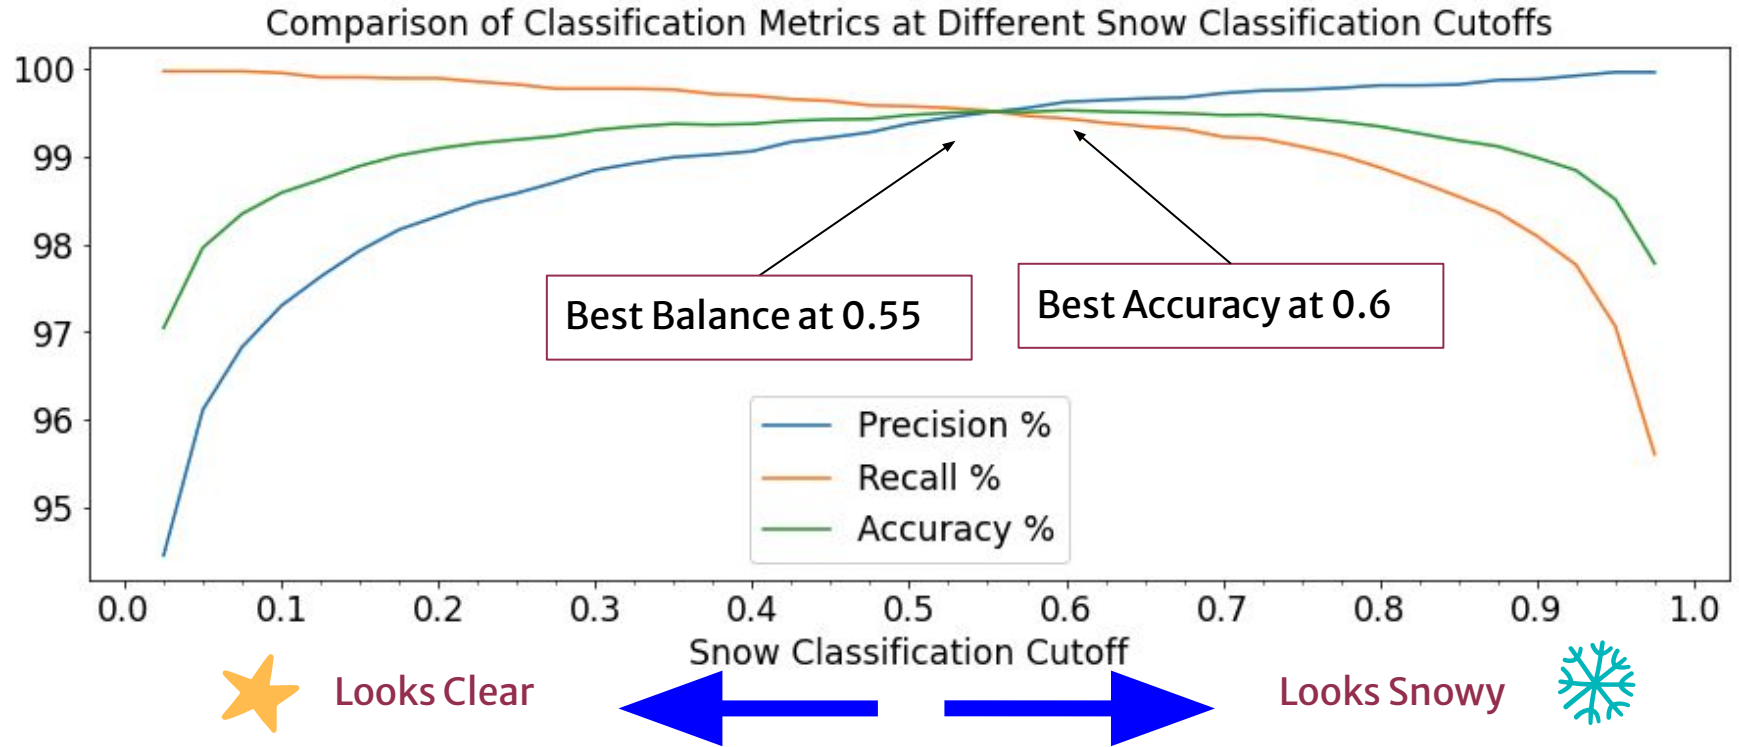
How tuneable is it?

Model 2


Convolution Kernel: 4x4

Max Pooling Kernel: 4x4


Training Passes: 10



Model 2



Convolution Kernel: 4x4



Max Pooling Kernel: 4x4

Training Passes: 10

Updated Metrics: Cutoff 0.55

Precision: If I say a picture has snow, how often am I right?

- 99.37% → 99.50%

Recall: Of the pictures with snow, how many did I identify?

- 99.57% → 99.52%

Accuracy: Overall, how many of my predictions were correct?

- 99.47% → 99.51%

This model is still highly tunable.

Streamlit

Let's see the model in action.



Conclusions

Takeaways

- The models presented are very capable of identifying falling snow in images.
- The models can be easily tuned for greater precision or recall depending on the application.
- The models are small, less than 1 MB for model2, making them easily deployable.

Future Directions

- Test on more real snow images.
- Analyze failure cases.
- Further training and tuning possible, but likely unnecessary.
 - Optimize for speed.
- Identify snow in live video or webcam. May require more sophisticated algorithm (YOLO).
- Use a GAN neural net to remove snow from images.
 - Original use for this data.





Thanks!

Do you have any questions?

CREDITS: This presentation template was created by **Slidesgo**, including icons by **Flaticon**, and infographics & images by **Freepik**