



# Automatic Snowfall Detection

Machine Vision for Identification of Snowfall in Images using a Convolutional Neural Network



# The Problem

**Can a neural network be trained to automatically recognize falling snow via machine vision?**

**U.S. Department of Transportation:  
~ 4% of vehicle crashes are related to snow**

**Localized snow detection could be useful for  
transportation and emergency personnel.**

**Navigation apps could inform users of falling snow.**

**On a separate note, ski resorts could easily monitor  
new snowfall on each run.**

# Presentation Overview



## 01.

### The Data

- Data source
- Overview

## 02.

### Preprocessing

- Preparing the data for modelling

## 03.

### The Models

- Structure
- Performance
- Streamlit app

## 04.

### Conclusions

- Takeaways
- Future Directions

# The Data

## DesnowNet: Context-Aware Deep Network for Snow Removal

Liu, Yun-Fu and Jaw, Da-Wei and Huang, Shih-Chia and Hwang, Jenq-Neng

- 50,000 images without snow
- A copy of each image with falling snow artificially added
- All RGB-encoded .jpg
- Various resolutions
- Largest dimension for any image = 640 pixels

Kindly provided for download on their website:

<https://sites.google.com/view/yunfuliu/desnownet>





## Challenges of this Dataset

1. 100,000 images is far too large for working memory
  - a. How can we shuffle and split?
  - b. How can we feed so many images to the model?
2. Images must be paired with a numerical label
  - a. 1 for snow
  - b. 0 for clear
3. Long time to train
4. Images are different sizes

# Preparing the Images for Modelling

1. 100,000 images is far too large for working memory
  - a. Start with just the file paths
  - b. Shuffle and split before reading in image data
  - c. Read in the images in batches

2. Images must be paired with a numerical label

3. Long time to train

Use a GPU →



4. Images are different sizes

Get the labels from the folder name



→ 0



→ 1

Buffer




# Ready for the Neural Net



1 for snow  
0 for clear

Buffers make every  
image 640x640

# The Models



Convolutional Layer: Groups pixels near each other, helping identify shapes.



Max Pooling Layer: Pixels near each other all given the value of the largest.



Dense Layer: Look for patterns. Mostly a black box.




# Model 1

Convolution Kernel: 5x5


Max Pooling Kernel: 3x3

Training Passes: 5


Training and Validation Accuracy at the End of Each Training Epoch



## Model 1



Convolution Kernel: 5x5



Max Pooling Kernel: 3x3

Training Passes: 5

## Metrics

Precision: If I say a picture has snow, how often am I right?

- 99.26%

Recall: Of the pictures with snow, how many did I identify?

- 97.80%

Accuracy: Overall, how many of my predictions were correct?

- 98.53%

I don't want to miss snowy pictures.

Can I improve the Recall?

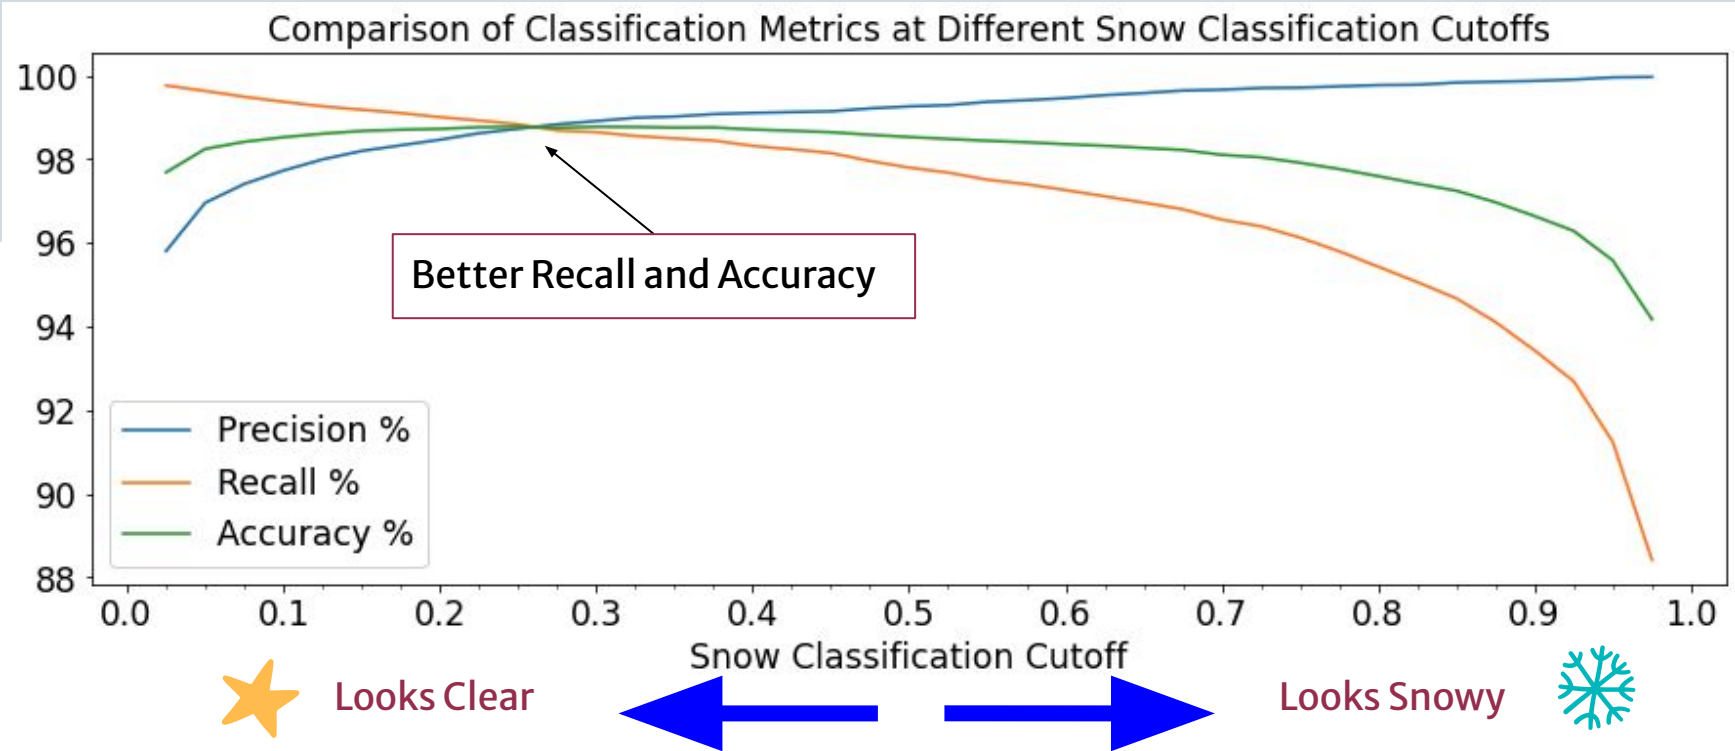
And keep high accuracy?

# Model 1


Convolution Kernel: 5x5

Max Pooling Kernel: 3x3


Training Passes: 5



## Model 1



Convolution Kernel: 5x5



Max Pooling Kernel: 3x3

Training Passes: 5

## Updated Metrics

Precision: If I say a picture has snow, how often am I right?

- 99.26% → 98.73%

Recall: Of the pictures with snow, how many did I identify?

- 97.80% → 98.84%

Accuracy: Overall, how many of my predictions were correct?

- 98.53% → 98.78%

Cutoff can be adjusted significantly without much accuracy loss.

Model is highly tunable.


## Model 2

Convolution Kernel: 4x4


Max Pooling Kernel: 4x4

Training Passes: 10


Training and Validation Accuracy at the End of Each Training Epoch



## Model 2



Convolution Kernel: 4x4



Max Pooling Kernel: 4x4

Training Passes: 10

## Metrics

Precision: If I say a picture has snow, how often am I right?

- 99.37%

Recall: Of the pictures with snow, how many did I identify?

- 99.57%

Accuracy: Overall, how many of my predictions were correct?

- 99.47%

This model is already very  
balanced.

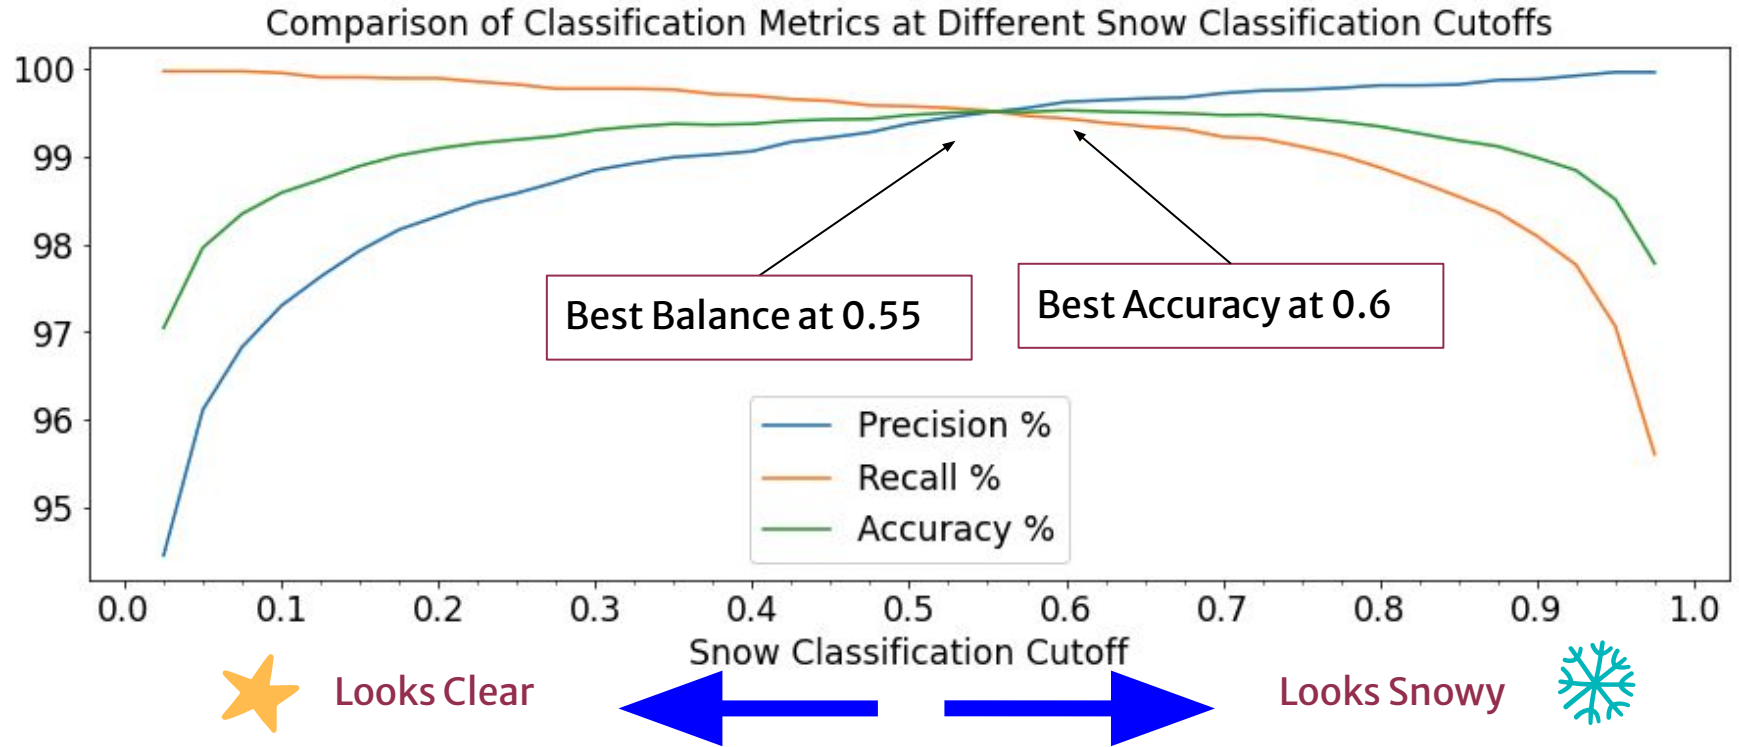
How tuneable is it?

## Model 2


Convolution Kernel: 4x4

Max Pooling Kernel: 4x4


Training Passes: 10



## Model 2



Convolution Kernel: 4x4



Max Pooling Kernel: 4x4

Training Passes: 10

## Updated Metrics: Cutoff 0.55

Precision: If I say a picture has snow, how often am I right?

- 99.37% → 99.50%

Recall: Of the pictures with snow, how many did I identify?

- 99.57% → 99.52%

Accuracy: Overall, how many of my predictions were correct?

- 99.47% → 99.51%

This model is still highly tunable.



**Streamlit**

Let's see the model in action.



# Conclusions

## Takeaways

- The models presented are very capable of identifying falling snow in images.
- The models can be easily tuned for greater precision or recall depending on the application.
- The models are small, less than 1 MB for model2, making them easily deployable.

## Future Directions

- Test on more real snow images.
- Analyze failure cases.
- Further training and tuning possible, but likely unnecessary.
  - Optimize for speed.
- Identify snow in live video or webcam. May require more sophisticated algorithm (YOLO).
- Use a GAN neural net to remove snow from images.
  - Original use for this data.





# Thanks!

Do you have any questions?

CREDITS: This presentation template was created by **Slidesgo**, including icons by **Flaticon**, and infographics & images by **Freepik**